

STSCI 4740 Final Project Group 10 – Appendix (R Output)

```
> setwd("C:/Users/14014/Documents/Cornell_Fall_2019/STS
CI_4740/STSCI_4740_FinalProject")
> divorce <- read.csv('divorce.csv', header=TRUE)
> #fix(divorce)
> attach(divorce)
The following object is masked _by_ .GlobalEnv:

      Class

The following objects are masked from divorce (pos = 3)
:

Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9,
      Class

The following objects are masked from divorce (pos = 4)
:

Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9,
      Class

The following objects are masked from divorce (pos = 14
):

Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9,
      Class

> #install.packages("Hmisc")
```

```

> #library(Hmisc)
> #describe(divorce)
> Class <- factor(Class)
> is.factor(Class)
[1] TRUE
> ### Quick correlation matrix output to see which parameters are good
> p <- 54
> k <- 5
> folds <- sample (1:k,nrow(divorce),replace=TRUE)
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+
+   corr_matrix <- matrix(c(rep(1,p),rep(1,p)), ncol=2)
+   colnames(corr_matrix) <- c("Atr","Correlation")
+
+   for (i in 1:54){
+     x <- c("Atr", i)
+     y <- paste(x, collapse= "")
+     correl = cor(train[ , i], train[ , 55])
+     corr_matrix[i,2] <- correl
+     corr_matrix[i,1] <- y
+   }
+   corr_matrix_sort <- corr_matrix[order(corr_matrix[,2]),]
+   print(tail(corr_matrix_sort))
+ }

```

	Atr	Correlation
[49,]	"Atr9"	"0.918764421686489"
[50,]	"Atr11"	"0.921040320729143"
[51,]	"Atr18"	"0.928324737691998"
[52,]	"Atr40"	"0.93176697034468"
[53,]	"Atr19"	"0.938507619834182"
[54,]	"Atr17"	"0.941020734208694"
	Atr	Correlation
[49,]	"Atr38"	"0.908224823593169"
[50,]	"Atr11"	"0.91909019979494"
[51,]	"Atr18"	"0.920965126287096"
[52,]	"Atr17"	"0.921998332657598"
[53,]	"Atr19"	"0.922263641646883"
[54,]	"Atr40"	"0.948874474286447"
	Atr	Correlation
[49,]	"Atr18"	"0.919625583000681"
[50,]	"Atr11"	"0.921042759515559"
[51,]	"Atr9"	"0.924236721302092"
[52,]	"Atr17"	"0.924903246391796"
[53,]	"Atr40"	"0.934627647404616"
[54,]	"Atr19"	"0.936738580233108"
	Atr	Correlation
[49,]	"Atr15"	"0.90485644723485"
[50,]	"Atr11"	"0.911740217681647"
[51,]	"Atr19"	"0.918678103080277"
[52,]	"Atr18"	"0.921457252524313"
[53,]	"Atr17"	"0.928342822102871"
[54,]	"Atr40"	"0.932726086069826"

```

      Atr      Correlation
[49,] "Atr41" "0.915393046745463"
[50,] "Atr11" "0.919512598526148"
[51,] "Atr18" "0.926739097887571"
[52,] "Atr19" "0.927495398189489"
[53,] "Atr17" "0.931827344566063"
[54,] "Atr40" "0.945048764279963"
> plot(Atr40, Class, main="Scatterplot of Atr40 vs. Cla
ss",
+       xlab="Atr40", ylab="Class", pch=19)
> set.seed(1)
> # linear regression #
> #1) Perform 5-fold CV to get mse
> lin.mod.mse_vector = rep(0,k)
> for(j in 1:k){
+   set.seed(1)
+
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   lin.mod = lm(Class~., data=train)
+   lin.mod_predict = predict(lin.mod, newdata = test)
+
+   error = mean((test$Class - lin.mod_predict)^2)
+   lin.mod.mse_vector[j]=error
+ }
> lin.mod.mse_vector
[1] 0.056688434 0.031478861 0.046281356 0.006454491
[5] 0.043046155
> mean(lin.mod.mse_vector)
[1] 0.03678986
> # Linear regression - all parameters / classification
threshold
> lin.mod.class.mse_vector = rep(0,k)
> for(j in 1:k){
+   set.seed(1)
+
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   lin.mod = lm(Class~., data=train)
+   lin.mod_predict = predict(lin.mod, newdata = test)
+   for(i in 1:length(lin.mod_predict)){
+     if(lin.mod_predict[i]>=0.5){
+       lin.mod_predict[i]=1
+     } else{
+       lin.mod_predict[i]=0
+     }
+   }
+
+   error = mean((test$Class - lin.mod_predict)^2)
+   lin.mod.class.mse_vector[j]=error
+ }
> lin.mod.class.mse_vector
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303

```

```

> mean(lin.mod.class.mse_vector)
[1] 0.02276863
> # Basic polynomial linear regression with value of 2
- all parameters
> log.mod.mse_vector = rep(0,k)
> for(j in 1:k){
+   set.seed(1)
+
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+
+   xnam=paste("Atr",2:54,sep="")
+   fmla=as.formula(paste("Class~poly(Atr1+",paste(xnam
, collapse="+"),paste(",2,raw=T)"))))
+   log.mod=lm(fmla,data=train) #use raw polynomial x,
x^2,x^3,...
+   log.mod_predict = predict(log.mod, newdata = test)
+   error = mean((test$Class - log.mod_predict)^2)
+   log.mod.mse_vector[j]=error
+ }
> log.mod.mse_vector
[1] 0.020820286 0.003320101 0.022064738 0.008646856
[5] 0.023494827
> mean(log.mod.mse_vector)
[1] 0.01566936
> # Polynomial - value of 2 - all parameters / include
classification threshold
> log.mod.class.mse_vector = rep(0,k)
> for(j in 1:k){
+   set.seed(1)
+
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+
+   xnam=paste("Atr",2:54,sep="")
+   fmla=as.formula(paste("Class~poly(Atr1+",paste(xnam
, collapse="+"),paste(",2,raw=T)"))))
+   log.mod=lm(fmla,data=train)
+   log.mod_predict = predict(log.mod, newdata = test)
+   for(i in 1:length(log.mod_predict)){
+     if(log.mod_predict[i]>=0.5){
+       log.mod_predict[i]=1
+     } else{
+       log.mod_predict[i]=0
+     }
+   }
+   error = mean((test$Class - log.mod_predict)^2)
+   log.mod.class.mse_vector[j]=error
+ }
> log.mod.class.mse_vector
[1] 0.02564103 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> mean(log.mod.class.mse_vector)
[1] 0.01764042
> # polynomial logistic regression #
> log.mod.mse_vector1 = rep(0,k)

```

```

> for(j in 1:k){
+   set.seed(1)
+
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+
+   xnam=paste("Atr",2:54,sep="")
+   fmla=as.formula(paste("Class~poly(Atr1+",paste(xnam
+ , collapse="+"),paste(",2,raw=T)"))))
+   log.mod=lm(fmla,data=train,family=binomial) #only difference: family=binomial
+   log.mod_predict = predict(log.mod, newdata = test)
+   error = mean((test$Class - log.mod_predict)^2)
+   log.mod.mse_vector1[j]=error
+ }
Warning messages:
1: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
  extra argument 'family' will be disregarded
2: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
  extra argument 'family' will be disregarded
3: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
  extra argument 'family' will be disregarded
4: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
  extra argument 'family' will be disregarded
5: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
  extra argument 'family' will be disregarded
> log.mod.mse_vector1
[1] 0.020820286 0.003320101 0.022064738 0.008646856
[5] 0.023494827
> mean(log.mod.mse_vector1)
[1] 0.01566936
> ### Perform basic LDA with all parameters
> library(MASS)
> lda.error=rep(0,k)
> for (j in 1:k){
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+   lda.fit=lda(as.factor(Class)~., data=train)
+   lda.fit
+
+   lda.pred=predict(lda.fit, test)
+   lda.class=lda.pred$class
+   #print(table(lda.class,class.test))
+   mean(lda.class==class.test)
+
+   topleft <- table(lda.class,class.test)[1]
+   bottomleft <- table(lda.class,class.test)[2]
+   topright <- table(lda.class,class.test)[3]
+   bottomright <- table(lda.class,class.test)[4]
+

```

```

+   lda.error[j]= (topright + bottomleft)/(topleft + bottomright + topright + bottomleft)
+ }
> lda.error
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> mean(lda.error)
[1] 0.02276863
> ### Perform CV for k in KNN with all parameters
> library(class)
> knn.error=rep(0,k)
> for (i in 1:k){
+   train = divorce[folds!=i,]
+   test = divorce[folds==i,]
+   class.test = Class[folds==i]
+   train.X = cbind(train[,c(1:54)])
+   test.X = cbind(test[,c(1:54)])
+   class.train = Class[folds!=i]
+
+   k.error=rep(0,100)
+
+   for(j in 1:100){
+     knn.pred=knn(train.X,test.X,class.train,k=j)
+     knn.class <- knn.pred
+     #print(table(knn.class,class.test))
+
+     topleft <- table(knn.class,class.test)[1]
+     bottomleft <- table(knn.class,class.test)[2]
+     topright <- table(knn.class,class.test)[3]
+     bottomright <- table(knn.class,class.test)[4]
+
+     error_rate <- (topright + bottomleft)/(topleft + bottomright + topright + bottomleft)
+     k.error[j] <- error_rate
+   }
+   k.error
+   min <- which.min(k.error)
+
+   knn.pred=knn(train.X,test.X,class.train,k=min)
+   knn.class <- knn.pred
+   table(knn.class,class.test)
+
+   error_rate <- (topright + bottomleft)/(topleft + bottomright + topright + bottomleft)
+   knn.error[i] <- error_rate
+ }
> knn.error
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> mean(knn.error)
[1] 0.02276863

>

> ### Basic Decision Tree (all 54 parameters)

```

```

> library(tree)
> tree.error=rep(0,k)
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+
+   tree.divorce=tree(as.factor(Class)~.,train)
+   plot(tree.divorce)
+   text(tree.divorce)
+   tree.pred=predict(tree.divorce,test,type="class")
+   tree.table=table(tree.pred,class.test)
+
+   topleft <- tree.table[1]
+   bottomleft <- tree.table[2]
+   topright <- tree.table[3]
+   bottomright <- tree.table[4]
+
+   tree.error[j]= (topright + bottomleft)/(topleft + b
ottomright + topright + bottomleft)
+ }
> tree.error
[1] 0.05128205 0.00000000 0.03225806 0.03030303
[5] 0.03030303
> mean(tree.error)
[1] 0.02882924
> ### Re-run LDA with only top 5 parameters
> lda5.error=rep(0,k)
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+   lda.fit=lda(as.factor(Class)~Atr40+Atr17+Atr19+Atr1
8+Atr11, data=train)
+   lda.fit
+
+   lda.pred=predict(lda.fit, test)
+   lda.class=lda.pred$class
+   table(lda.class,class.test)
+   mean(lda.class==class.test)
+
+   topleft <- table(lda.class,class.test)[1]
+   bottomleft <- table(lda.class,class.test)[2]
+   topright <- table(lda.class,class.test)[3]
+   bottomright <- table(lda.class,class.test)[4]
+
+   lda5.error[j]= (topright + bottomleft)/(topleft + b
ottomright + topright + bottomleft)
+ }
> lda5.error
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> mean(lda5.error)
[1] 0.02276863

```

```

> ### Re-run QDA with only top 5 parameters - should work now
> qda5.error=rep(0,k)
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+
+   qda.fit=qda(as.factor(Class)~Atr40+Atr17+Atr19+Atr18+Atr11, data=train)
+   qda.class=predict(qda.fit,test)$class
+   print(table(qda.class,class.test))
+   mean(qda.class==class.test)
+
+   topleft <- table(qda.class,class.test)[1]
+   bottomleft <- table(qda.class,class.test)[2]
+   topright <- table(qda.class,class.test)[3]
+   bottomright <- table(qda.class,class.test)[4]
+
+   qda5.error[j]= (topright + bottomleft)/(topleft + bottomright + topright + bottomleft)
+ }

      class.test
qda.class  0  1
      0 20  1
      1  0 18
      class.test
qda.class  0  1
      0 13  0
      1  0 21
      class.test
qda.class  0  1
      0 18  1
      1  0 12
      class.test
qda.class  0  1
      0 18  0
      1  0 15
      class.test
qda.class  0  1
      0 17  1
      1  0 15
> qda5.error
[1] 0.02564103 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> mean(qda5.error)
[1] 0.01764042
> ### CV for k with KNN with only top 5 parameters
> knn5.error=rep(0,k)
> for (i in 1:k){
+   set.seed(1)
+   train = divorce[folds!=i,]
+   test = divorce[folds==i,]
+   class.test = Class[folds==i]
+   train.X = cbind(train[,c(11,17,18,19,40)])

```



```

+ test.X = cbind(test[,c(11,17,18,19,40)])
+ class.train = Class[folds!=i]
+
+ k.error=rep(0,100)
+
+ for(j in 1:100){
+   knn.pred=knn(train.X,test.X,class.train,k=j)
+   knn.class <- knn.pred
+   table(knn.class,class.test)
+
+   topleft <- table(knn.class,class.test)[1]
+   bottomleft <- table(knn.class,class.test)[2]
+   topright <- table(knn.class,class.test)[3]
+   bottomright <- table(knn.class,class.test)[4]
+
+   error_rate <- (topright + bottomleft)/(topleft +
bottomright + topright + bottomleft)
+   k.error[j] <- error_rate
+ }
+ k.error
+ min <- which.min(k.error)
+ print(min)
+
+ knn.pred=knn(train.X,test.X,class.train,k=min)
+ knn.class <- knn.pred
+ table(knn.class,class.test)
+
+ error_rate <- (topright + bottomleft)/(topleft + bo
ttomright + topright + bottomleft)
+ knn5.error[i] <- error_rate
+
+ }
[1] 1
[1] 1
[1] 1
[1] 1
[1] 1
[1] 1
> knn5.error
[1] 0.48717949 0.61764706 0.03225806 0.00000000
[5] 0.48484848
> mean(knn5.error)
[1] 0.3243866
> ### Decision Tree with 5 parameters only
> tree5.error=rep(0,k)
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+
+   tree5.divorce=tree(as.factor(Class)~.,train)
+   plot(tree5.divorce)
+   text(tree5.divorce)
+   tree5.pred=predict(tree5.divorce,test,type="class")
+   tree5.table=table(tree5.pred,class.test)
+

```

```

+   topleft <- tree.table[1]
+   bottomleft <- tree.table[2]
+   topright <- tree.table[3]
+   bottomright <- tree.table[4]
+
+   tree5.error[j]= (topright + bottomleft)/(topleft +
bottomright + topright + bottomleft)
+ }
> tree5.error
[1] 0.03030303 0.03030303 0.03030303 0.03030303
[5] 0.03030303
> mean(tree5.error)
[1] 0.03030303
> xnam=paste("Atr",1:54,sep="")
> fmla=as.formula(paste("Class~",paste(xnam, collapse="
+")))
> fmla
Class ~ Atr1 + Atr2 + Atr3 + Atr4 + Atr5 + Atr6 + Atr7
+ Atr8 +
      Atr9 + Atr10 + Atr11 + Atr12 + Atr13 + Atr14 + Atr1
5 + Atr16 +
      Atr17 + Atr18 + Atr19 + Atr20 + Atr21 + Atr22 + Atr
23 + Atr24 +
      Atr25 + Atr26 + Atr27 + Atr28 + Atr29 + Atr30 + Atr
31 + Atr32 +
      Atr33 + Atr34 + Atr35 + Atr36 + Atr37 + Atr38 + Atr
39 + Atr40 +
      Atr41 + Atr42 + Atr43 + Atr44 + Atr45 + Atr46 + Atr
47 + Atr48 +
      Atr49 + Atr50 + Atr51 + Atr52 + Atr53 + Atr54
> # First Step:
> #fitting the model without using any subset selection
s
> library(boot)
> set.seed(1)
> cv.logi_error1.5=rep(0,5)
> for (i in 1:5){
+   logi.fit=glm(Class~.,data=divorce,family='binomial'
)
+   cv.logi_error1.5[i]=cv.glm(divorce,logi.fit, K=5)$d
elta[1]
+ }
There were 50 or more warnings (use warnings() to see t
he first 50)
> cv.logi_error1.5
[1] 0.03449928 0.05066133 0.02453584 0.04774724
[5] 0.03648031
> mean(cv.logi_error1.5)
[1] 0.0387848
> cv.logiWC_error1.5=rep(0,5)
> for (i in 1:5){
+   logiWC.fit=glm(Class~., data=divorce, family='binom
ial',control=list(maxit=1000))
+   cv.logiWC_error1.5[i]=cv.glm(divorce,logiWC.fit, K=
5)$delta[1]
+ }

```

There were 30 warnings (use warnings() to see them)

```
> cv.logiWC_error1.5
[1] 0.03220952 0.03154547 0.02935895 0.02339867
[5] 0.04088782
> mean(cv.logiWC_error1.5)
[1] 0.03148009
> # ***** GAM Setup *****
> library(gam)
> newx=paste("s(Atr",1:54,',df=k)',sep="")
> fmla2=as.formula(paste('I(Class==1)~',paste(newx, col
lapse="+")))
> fmla2
I(Class == 1) ~ s(Atr1, df = k) + s(Atr2, df = k) + s(Atr3, df = k) +
s(Atr4, df = k) + s(Atr5, df = k) + s(Atr6, df = k)
+ s(Atr7,
df = k) + s(Atr8, df = k) + s(Atr9, df = k) + s(Atr
10, df = k) +
s(Atr11, df = k) + s(Atr12, df = k) + s(Atr13, df =
k) +
s(Atr14, df = k) + s(Atr15, df = k) + s(Atr16, df =
k) +
s(Atr17, df = k) + s(Atr18, df = k) + s(Atr19, df =
k) +
s(Atr20, df = k) + s(Atr21, df = k) + s(Atr22, df =
k) +
s(Atr23, df = k) + s(Atr24, df = k) + s(Atr25, df =
k) +
s(Atr26, df = k) + s(Atr27, df = k) + s(Atr28, df =
k) +
s(Atr29, df = k) + s(Atr30, df = k) + s(Atr31, df =
k) +
s(Atr32, df = k) + s(Atr33, df = k) + s(Atr34, df =
k) +
s(Atr35, df = k) + s(Atr36, df = k) + s(Atr37, df =
k) +
s(Atr38, df = k) + s(Atr39, df = k) + s(Atr40, df =
k) +
s(Atr41, df = k) + s(Atr42, df = k) + s(Atr43, df =
k) +
s(Atr44, df = k) + s(Atr45, df = k) + s(Atr46, df =
k) +
s(Atr47, df = k) + s(Atr48, df = k) + s(Atr49, df =
k) +
s(Atr50, df = k) + s(Atr51, df = k) + s(Atr52, df =
k) +
s(Atr53, df = k) + s(Atr54, df = k)
```

```
> gam.fit=gam(fmla2,family=binomial,data=divorce)
```

Warning messages:

```
1: In model.matrix.default(mt, mf, contrasts) :
non-list contrasts argument ignored
2: In gam(fmla2, family = binomial, data = divorce) :
Residual degrees of freedom are negative or zero. Th
is occurs when the sum of the parametric and nonparamet
ric degrees of freedom exceeds the number of observatio
```

ns. The model is probably too complex for the amount of data available.

```
> summary(gam.fit)
```

```
Call: gam(formula = fmla2, family = binomial, data = di  
vorce)
```

```
(Dispersion Parameter for binomial family taken to be 1  
)
```

```
Null Deviance: 235.6465 on 169 degrees of freedom  
Residual Deviance: 0 on -47 degrees of freedom  
AIC: 434.0001
```

```
Number of Local Scoring Iterations: 24
```

```
Anova for Parametric Effects
```

	Df	Sum Sq	Mean Sq
s(Atr1, df = k)	1	1.6181e-06	1.6181e-06
s(Atr2, df = k)	1	8.3070e-08	8.3070e-08
s(Atr3, df = k)	1	1.4720e-08	1.4720e-08
s(Atr4, df = k)	1	2.2880e-08	2.2880e-08
s(Atr5, df = k)	1	1.1392e-07	1.1392e-07
s(Atr6, df = k)	1	9.9140e-08	9.9140e-08
s(Atr7, df = k)	1	1.0320e-08	1.0320e-08
s(Atr8, df = k)	1	6.1670e-08	6.1670e-08
s(Atr9, df = k)	1	1.9430e-08	1.9430e-08
s(Atr10, df = k)	1	3.2000e-10	3.2000e-10
s(Atr11, df = k)	1	2.9860e-08	2.9860e-08
s(Atr12, df = k)	1	9.9100e-09	9.9100e-09
s(Atr13, df = k)	1	5.9200e-09	5.9200e-09
s(Atr14, df = k)	1	9.4000e-10	9.4000e-10
s(Atr15, df = k)	1	2.6900e-09	2.6900e-09
s(Atr16, df = k)	1	7.2000e-10	7.2000e-10
s(Atr17, df = k)	1	3.0200e-09	3.0200e-09
s(Atr18, df = k)	1	2.7700e-09	2.7700e-09
s(Atr19, df = k)	1	1.6000e-10	1.6000e-10
s(Atr20, df = k)	1	1.3000e-10	1.3000e-10
s(Atr21, df = k)	1	0.0000e+00	0.0000e+00
s(Atr22, df = k)	1	1.0750e-08	1.0750e-08
s(Atr23, df = k)	1	4.0200e-09	4.0200e-09
s(Atr24, df = k)	1	2.9100e-09	2.9100e-09
s(Atr25, df = k)	1	8.8600e-09	8.8600e-09
s(Atr26, df = k)	1	2.0880e-08	2.0880e-08
s(Atr27, df = k)	1	4.2200e-09	4.2200e-09
s(Atr28, df = k)	1	7.4000e-09	7.4000e-09
s(Atr29, df = k)	1	7.9000e-10	7.9000e-10
s(Atr30, df = k)	1	2.2200e-09	2.2200e-09
s(Atr31, df = k)	1	7.4500e-09	7.4500e-09
s(Atr32, df = k)	1	4.1000e-10	4.1000e-10
s(Atr33, df = k)	1	2.6000e-10	2.6000e-10
s(Atr34, df = k)	1	2.0100e-08	2.0100e-08
s(Atr35, df = k)	1	5.7200e-09	5.7200e-09
s(Atr36, df = k)	1	6.6300e-09	6.6300e-09
s(Atr37, df = k)	1	0.0000e+00	0.0000e+00
s(Atr38, df = k)	1	4.5000e-10	4.5000e-10

s(Atr39, df = k)	1	2.6400e-09	2.6400e-09
s(Atr40, df = k)	1	1.1890e-08	1.1890e-08
s(Atr41, df = k)	1	4.0000e-11	4.0000e-11
s(Atr42, df = k)	1	2.1000e-10	2.1000e-10
s(Atr43, df = k)	1	5.2000e-10	5.2000e-10
s(Atr44, df = k)	1	8.0000e-10	8.0000e-10
s(Atr45, df = k)	1	1.0000e-10	1.0000e-10
s(Atr46, df = k)	1	7.0000e-11	7.0000e-11
s(Atr47, df = k)	1	3.2400e-09	3.2400e-09
s(Atr48, df = k)	1	7.4000e-10	7.4000e-10
s(Atr49, df = k)	1	1.4800e-09	1.4800e-09
s(Atr50, df = k)	1	1.0000e-11	1.0000e-11
s(Atr51, df = k)	1	1.8100e-09	1.8100e-09
s(Atr52, df = k)	1	1.4950e-08	1.4950e-08
s(Atr53, df = k)	1	1.5000e-10	1.5000e-10
s(Atr54, df = k)	1	2.0000e-11	2.0000e-11
Residuals	-47	0.0000e+00	0.0000e+00

F value Pr(>F)

s(Atr1, df = k)	-2.9554e+08
s(Atr2, df = k)	-1.5173e+07
s(Atr3, df = k)	-2.6880e+06
s(Atr4, df = k)	-4.1780e+06
s(Atr5, df = k)	-2.0807e+07
s(Atr6, df = k)	-1.8107e+07
s(Atr7, df = k)	-1.8848e+06
s(Atr8, df = k)	-1.1264e+07
s(Atr9, df = k)	-3.5482e+06
s(Atr10, df = k)	-5.8766e+04
s(Atr11, df = k)	-5.4540e+06
s(Atr12, df = k)	-1.8096e+06
s(Atr13, df = k)	-1.0816e+06
s(Atr14, df = k)	-1.7119e+05
s(Atr15, df = k)	-4.9207e+05
s(Atr16, df = k)	-1.3084e+05
s(Atr17, df = k)	-5.5226e+05
s(Atr18, df = k)	-5.0630e+05
s(Atr19, df = k)	-2.9870e+04
s(Atr20, df = k)	-2.2940e+04
s(Atr21, df = k)	-5.3413e+02
s(Atr22, df = k)	-1.9637e+06
s(Atr23, df = k)	-7.3419e+05
s(Atr24, df = k)	-5.3062e+05
s(Atr25, df = k)	-1.6182e+06
s(Atr26, df = k)	-3.8128e+06
s(Atr27, df = k)	-7.7026e+05
s(Atr28, df = k)	-1.3508e+06
s(Atr29, df = k)	-1.4409e+05
s(Atr30, df = k)	-4.0557e+05
s(Atr31, df = k)	-1.3611e+06
s(Atr32, df = k)	-7.4995e+04
s(Atr33, df = k)	-4.6908e+04
s(Atr34, df = k)	-3.6717e+06
s(Atr35, df = k)	-1.0453e+06
s(Atr36, df = k)	-1.2111e+06
s(Atr37, df = k)	-8.4973e+02
s(Atr38, df = k)	-8.1957e+04

```

s(Atr39, df = k) -4.8309e+05
s(Atr40, df = k) -2.1717e+06
s(Atr41, df = k) -7.9474e+03
s(Atr42, df = k) -3.8962e+04
s(Atr43, df = k) -9.5392e+04
s(Atr44, df = k) -1.4645e+05
s(Atr45, df = k) -1.8223e+04
s(Atr46, df = k) -1.2564e+04
s(Atr47, df = k) -5.9165e+05
s(Atr48, df = k) -1.3523e+05
s(Atr49, df = k) -2.7010e+05
s(Atr50, df = k) -1.7300e+03
s(Atr51, df = k) -3.2995e+05
s(Atr52, df = k) -2.7301e+06
s(Atr53, df = k) -2.7463e+04
s(Atr54, df = k) -3.8505e+03
Residuals

```

Anova for Nonparametric Effects

	Npar	Df	Npar	Chisq	P(Chi)
(Intercept)					
s(Atr1, df = k)	3	1.0202e-09		1	
s(Atr2, df = k)	3	1.2342e-09		1	
s(Atr3, df = k)	3	3.1302e-09		1	
s(Atr4, df = k)	3	2.8132e-09		1	
s(Atr5, df = k)	3	3.1820e-10		1	
s(Atr6, df = k)	3	4.3560e-10		1	
s(Atr7, df = k)	3	3.3831e-09		1	
s(Atr8, df = k)	3	1.2689e-09		1	
s(Atr9, df = k)	3	4.3280e-09		1	
s(Atr10, df = k)	3	7.6840e-10		1	
s(Atr11, df = k)	3	1.5846e-09		1	
s(Atr12, df = k)	3	4.2030e-10		1	
s(Atr13, df = k)	3	7.8200e-11		1	
s(Atr14, df = k)	3	1.0869e-09		1	
s(Atr15, df = k)	3	9.0400e-10		1	
s(Atr16, df = k)	3	7.8420e-10		1	
s(Atr17, df = k)	3	2.8769e-09		1	
s(Atr18, df = k)	3	2.4578e-09		1	
s(Atr19, df = k)	3	5.7110e-10		1	
s(Atr20, df = k)	3	1.5749e-09		1	
s(Atr21, df = k)	3	4.0171e-09		1	
s(Atr22, df = k)	3	6.2650e-10		1	
s(Atr23, df = k)	3	1.3182e-09		1	
s(Atr24, df = k)	3	3.8708e-09		1	
s(Atr25, df = k)	3	4.2460e-10		1	
s(Atr26, df = k)	3	3.9860e-09		1	
s(Atr27, df = k)	3	9.9800e-11		1	
s(Atr28, df = k)	3	2.7089e-09		1	
s(Atr29, df = k)	3	5.9623e-09		1	
s(Atr30, df = k)	3	1.3760e-10		1	
s(Atr31, df = k)	3	3.5530e-09		1	
s(Atr32, df = k)	3	5.4540e-10		1	
s(Atr33, df = k)	3	1.6051e-09		1	
s(Atr34, df = k)	3	5.1030e-10		1	
s(Atr35, df = k)	3	6.3010e-10		1	

```

s(Atr36, df = k)      3 2.7341e-09      1
s(Atr37, df = k)      3 4.1581e-09      1
s(Atr38, df = k)      3 3.6130e-10      1
s(Atr39, df = k)      3 7.3450e-10      1
s(Atr40, df = k)      3 3.4084e-09      1
s(Atr41, df = k)      3 4.6420e-10      1
s(Atr42, df = k)      3 4.9250e-10      1
s(Atr43, df = k)      3 9.7910e-10      1
s(Atr44, df = k)      3 1.2413e-09      1
s(Atr45, df = k)      3 2.8395e-09      1
s(Atr46, df = k)      3 2.4580e-10      1
s(Atr47, df = k)      3 3.4914e-09      1
s(Atr48, df = k)      3 9.4170e-10      1
s(Atr49, df = k)      3 2.5830e-10      1
s(Atr50, df = k)      3 5.0480e-10      1
s(Atr51, df = k)      3 4.3760e-10      1
s(Atr52, df = k)      3 3.8671e-09      1
s(Atr53, df = k)      3 1.9148e-09      1
s(Atr54, df = k)      3 2.6122e-09      1
Warning messages:
1: In pf(f, df, dfr, lower.tail = FALSE) : NaNs produce
d
2: In summary.Gam(gam.fit) :
  Residual degrees of freedom are negative or zero. Th
is occurs when the sum of the parametric and nonparamet
ric degrees of freedom exceeds the number of observatio
ns. The model is probably too complex for the amount o
f data available.
> gam.error=rep(0,5)
> # ***** GAM with Smoothing Splines with degr
ee of freedom=5 *****
> for (j in 1:5){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds==j,]
+   class.test = Class[folds==j]
+
+   gam.fit=gam(fmla2,family=binomial,data=train)
+   gam.pred=predict(gam.fit,test,type='response')
+   conf_gam= table(gam.pred>.5,class.test)
+
+   print(conf_gam)
+
+   #Apply it to the test data and get the confusion ma
trix and error rate.
+   gam.error[j]=1-sum(diag(conf_gam))/sum(conf_gam)
+ }

      class.test
      0  1
FALSE 20  2
TRUE   0 17
      class.test
      0  1
FALSE 13  1
TRUE   0 20
      class.test

```

```

      0  1
FALSE 18  1
TRUE   0 12
      class.test
      0  1
FALSE 18  0
TRUE   0 15
      class.test
      0  1
FALSE 17  1
TRUE   0 15
Warning messages:
1: In model.matrix.default(mt, mf, contrasts) :
  non-list contrasts argument ignored
2: In gam(fmla2, family = binomial, data = train) :
  Residual degrees of freedom are negative or zero. This occurs when the sum of the parametric and nonparametric degrees of freedom exceeds the number of observations. The model is probably too complex for the amount of data available.
3: In model.matrix.default(mt, mf, contrasts) :
  non-list contrasts argument ignored
4: In gam(fmla2, family = binomial, data = train) :
  Residual degrees of freedom are negative or zero. This occurs when the sum of the parametric and nonparametric degrees of freedom exceeds the number of observations. The model is probably too complex for the amount of data available.
5: In model.matrix.default(mt, mf, contrasts) :
  non-list contrasts argument ignored
6: In gam(fmla2, family = binomial, data = train) :
  Residual degrees of freedom are negative or zero. This occurs when the sum of the parametric and nonparametric degrees of freedom exceeds the number of observations. The model is probably too complex for the amount of data available.
7: In model.matrix.default(mt, mf, contrasts) :
  non-list contrasts argument ignored
8: In gam(fmla2, family = binomial, data = train) :
  Residual degrees of freedom are negative or zero. This occurs when the sum of the parametric and nonparametric degrees of freedom exceeds the number of observations. The model is probably too complex for the amount of data available.
9: In model.matrix.default(mt, mf, contrasts) :
  non-list contrasts argument ignored
10: In gam(fmla2, family = binomial, data = train) :
  Residual degrees of freedom are negative or zero. This occurs when the sum of the parametric and nonparametric degrees of freedom exceeds the number of observations. The model is probably too complex for the amount of data available.
> gam.error
[1] 0.05128205 0.02941176 0.03225806 0.00000000
[5] 0.03030303
> mean(gam.error)

```



```

[1] 0.02865098
> # ***** GAM with natural splines with degree of freedom = 5 *****
> library(splines)
> newx_ns=paste("ns(Atr",1:54,',df=k)',sep="")
> fmla_ns=as.formula(paste('Class~',paste(newx_ns, collapse="+")))
> fmla_ns
Class ~ ns(Atr1, df = k) + ns(Atr2, df = k) + ns(Atr3, df = k) +
      ns(Atr4, df = k) + ns(Atr5, df = k) + ns(Atr6, df = k) +
      ns(Atr7, df = k) + ns(Atr8, df = k) + ns(Atr9, df = k) +
      ns(Atr10, df = k) + ns(Atr11, df = k) + ns(Atr12, df = k) +
      ns(Atr13, df = k) + ns(Atr14, df = k) + ns(Atr15, df = k) +
      ns(Atr16, df = k) + ns(Atr17, df = k) + ns(Atr18, df = k) +
      ns(Atr19, df = k) + ns(Atr20, df = k) + ns(Atr21, df = k) +
      ns(Atr22, df = k) + ns(Atr23, df = k) + ns(Atr24, df = k) +
      ns(Atr25, df = k) + ns(Atr26, df = k) + ns(Atr27, df = k) +
      ns(Atr28, df = k) + ns(Atr29, df = k) + ns(Atr30, df = k) +
      ns(Atr31, df = k) + ns(Atr32, df = k) + ns(Atr33, df = k) +
      ns(Atr34, df = k) + ns(Atr35, df = k) + ns(Atr36, df = k) +
      ns(Atr37, df = k) + ns(Atr38, df = k) + ns(Atr39, df = k) +
      ns(Atr40, df = k) + ns(Atr41, df = k) + ns(Atr42, df = k) +
      ns(Atr43, df = k) + ns(Atr44, df = k) + ns(Atr45, df = k) +
      ns(Atr46, df = k) + ns(Atr47, df = k) + ns(Atr48, df = k) +
      ns(Atr49, df = k) + ns(Atr50, df = k) + ns(Atr51, df = k) +
      ns(Atr52, df = k) + ns(Atr53, df = k) + ns(Atr54, df = k)
> ## Code for best subset selection within k-fold for linear regression ##
> set.seed(1)

>

> val <- 8 #max preds for bestsubset
> best.lin.cv.errors = matrix (NA,k,val, dimnames =list (NULL , paste (1:val) ))
> #install.packages("leaps")
> library(leaps)
> # predict function for best subset selection #

```

```

> predict.regsubsets = function(object,newdata,id,...){
+   form=as.formula (object$call [[2]])
+   mat=model.matrix(form ,newdata )
+   coefi=coef(object ,id=id)
+   xvars=names(coefi)
+   mat[,xvars]%%coefi
+ }
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   best.fit=regsubsets(Class~.,data=train,nvmax=val,re
ally.big=T)
+   for(i in 1:val){
+     pred=predict(best.fit,test,id=i)
+     for(m in 1:length(pred)){
+       if(pred[m]>=0.5){
+         pred[m]=1
+       } else{
+         pred[m]=0
+       }
+     }
+     best.lin.cv.errors[j,i]= mean((test$Class-pred)^2
)
+   }
+ }
> best.lin.cv.errors
      1      2      3      4
[1,] 0.05128205 0.05128205 0.02564103 0.02564103
[2,] 0.05882353 0.02941176 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.00000000 0.03225806
[4,] 0.03030303 0.00000000 0.00000000 0.00000000
[5,] 0.06060606 0.03030303 0.03030303 0.03030303
      5      6      7      8
[1,] 0.02564103 0.02564103 0.02564103 0.02564103
[2,] 0.02941176 0.00000000 0.02941176 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
> mean.best.lin.cv.errors=apply(best.lin.cv.errors ,2,
mean)
> mean.best.lin.cv.errors
      1      2      3      4
0.04665455 0.02865098 0.01118881 0.01764042
      5      6      7      8
0.02352278 0.01764042 0.02352278 0.01764042
> x = which.min(mean.best.lin.cv.errors)
> reg.best=regsubsets(Class~.,data=divorce , nvmax=val,
really.big=T)
> coef(reg.best,x)
(Intercept)      Atr6      Atr18      Atr40
-0.04022481  0.07411876  0.14921847  0.13498984
> ## Code for forward selection within k-fold for linea
r regression ##
> set.seed(1)

```

```

> val <- 54 #max preds for bestsubset
> fwd.lin.cv.errors = matrix (NA,k,val, dimnames =list(
NULL , paste (1:val) ))
> #install.packages("leaps")
> library(leaps)
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   regfit.fwd.fit=regsubsets(Class~.,data=train,nvmax=
val,method="forward")
+   for(i in 1:val){
+     pred=predict(regfit.fwd.fit,test,id=i)
+
+     for(m in 1:length(pred)){
+       if(pred[m]>=0.5){
+         pred[m]=1.0
+       } else{
+         pred[m]=0.0
+       }
+     }
+
+     fwd.lin.cv.errors[j,i]= mean((test$Class-pred)^2)
+   }
+ }
> fwd.lin.cv.errors

```

	1	2	3	4
[1,]	0.05128205	0.05128205	0.05128205	0.02564103
[2,]	0.05882353	0.02941176	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.03030303	0.00000000	0.00000000	0.00000000
[5,]	0.06060606	0.03030303	0.03030303	0.03030303
	5	6	7	8
[1,]	0.02564103	0.02564103	0.02564103	0.02564103
[2,]	0.02941176	0.00000000	0.02941176	0.02941176
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	9	10	11	12
[1,]	0.02564103	0.02564103	0.02564103	0.02564103
[2,]	0.02941176	0.02941176	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	13	14	15	16
[1,]	0.02564103	0.02564103	0.02564103	0.02564103
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	17	18	19	20
[1,]	0.02564103	0.02564103	0.02564103	0.02564103
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000

[illegible]

```

> mean.fwd.lin.cv.errors=apply(fwd.lin.cv.errors ,2, me
an)
> mean.fwd.lin.cv.errors
      1      2      3      4
0.04665455 0.02865098 0.02276863 0.01764042
      5      6      7      8
0.02352278 0.01764042 0.02352278 0.02352278
      9     10     11     12
0.02352278 0.02352278 0.01764042 0.01764042
     13     14     15     16
0.01764042 0.01764042 0.01764042 0.01764042
     17     18     19     20
0.01764042 0.01764042 0.01764042 0.01764042
     21     22     23     24
0.01764042 0.01764042 0.01764042 0.01764042
     25     26     27     28
0.01764042 0.01764042 0.01764042 0.01764042
     29     30     31     32
0.01764042 0.01764042 0.01764042 0.01764042
     33     34     35     36
0.01764042 0.01764042 0.01764042 0.02276863
     37     38     39     40
0.01764042 0.02276863 0.02276863 0.02276863
     41     42     43     44
0.02276863 0.02276863 0.02276863 0.02276863
     45     46     47     48
0.02276863 0.02276863 0.02276863 0.02276863
     49     50     51     52
0.02276863 0.02276863 0.02276863 0.02276863
     53     54
0.02276863 0.02276863
> mean(mean.fwd.lin.cv.errors)
[1] 0.02073065
> x = which.min(mean.fwd.lin.cv.errors)
> reg.best=regsubsets(Class~.,data=divorce , nvmax=val,
really.big=T,method="forward")
> coef(reg.best,x)
(Intercept)      Atr6      Atr18      Atr29
-0.03992621  0.07017789  0.12168790  0.04892915
      Atr40
0.11965844
> ## Code for backward selection within k-fold for line
ar regression ##
> set.seed(1)
> val <- 54 #max preds for bestsubset
> bwd.lin.cv.errors = matrix (NA,k,val, dimnames =list(
NULL , paste (1:val) ))
> #install.packages("leaps")
> library(leaps)
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   regfit.fwd.fit=regsubsets(Class~.,data=train,nvmax=
val,method="backward")

```

```

+   for(i in 1:val){
+     pred=predict(regfit.fwd.fit,test,id=i)
+     for(m in 1:length(pred)){
+       if(pred[m]>=0.5){
+         pred[m]=1
+       } else{
+         pred[m]=0
+       }
+     }
+     bwd.lin.cv.errors[j,i]= mean((test$Class-pred)^2)
+   }
+ }
> bwd.lin.cv.errors

```

	1	2	3	4
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.05882353	0.02941176	0.00000000	0.00000000
[3,]	0.00000000	0.03225806	0.03225806	0.03225806
[4,]	0.03030303	0.00000000	0.00000000	0.00000000
[5,]	0.06060606	0.03030303	0.03030303	0.03030303
	5	6	7	8
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	9	10	11	12
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	13	14	15	16
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	17	18	19	20
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	21	22	23	24
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	25	26	27	28
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303
	29	30	31	32
[1,]	0.05128205	0.05128205	0.05128205	0.05128205
[2,]	0.00000000	0.00000000	0.00000000	0.00000000
[3,]	0.03225806	0.03225806	0.03225806	0.03225806
[4,]	0.00000000	0.00000000	0.00000000	0.00000000
[5,]	0.03030303	0.03030303	0.03030303	0.03030303

```

[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      33      34      35      36
[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      37      38      39      40
[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      41      42      43      44
[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      45      46      47      48
[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      49      50      51      52
[1,] 0.05128205 0.05128205 0.05128205 0.05128205
[2,] 0.00000000 0.00000000 0.00000000 0.00000000
[3,] 0.03225806 0.03225806 0.03225806 0.03225806
[4,] 0.00000000 0.00000000 0.00000000 0.00000000
[5,] 0.03030303 0.03030303 0.03030303 0.03030303
      53      54
[1,] 0.05128205 0.05128205
[2,] 0.00000000 0.00000000
[3,] 0.03225806 0.03225806
[4,] 0.00000000 0.00000000
[5,] 0.03030303 0.03030303
> mean.bwd.lin.cv.errors=apply(bwd.lin.cv.errors ,2, me
an)
> mean.bwd.lin.cv.errors
      1      2      3      4
0.04020293 0.02865098 0.02276863 0.02276863
      5      6      7      8
0.02276863 0.02276863 0.02276863 0.02276863
      9     10     11     12
0.02276863 0.02276863 0.02276863 0.02276863
     13     14     15     16
0.02276863 0.02276863 0.02276863 0.02276863
     17     18     19     20
0.02276863 0.02276863 0.02276863 0.02276863
     21     22     23     24
0.02276863 0.02276863 0.02276863 0.02276863

```

```

      25      26      27      28
0.02276863 0.02276863 0.02276863 0.02276863
      29      30      31      32
0.02276863 0.02276863 0.02276863 0.02276863
      33      34      35      36
0.02276863 0.02276863 0.02276863 0.02276863
      37      38      39      40
0.02276863 0.02276863 0.02276863 0.02276863
      41      42      43      44
0.02276863 0.02276863 0.02276863 0.02276863
      45      46      47      48
0.02276863 0.02276863 0.02276863 0.02276863
      49      50      51      52
0.02276863 0.02276863 0.02276863 0.02276863
      53      54
0.02276863 0.02276863
> mean(mean.bwd.lin.cv.errors)
[1] 0.02320042
> x = which.min(mean.bwd.lin.cv.errors)
> reg.best=regsubsets(Class~.,data=divorce , nvmax=val,
really.big=T,method="backward")
> coef(reg.best,x)
(Intercept)      Atr17      Atr26      Atr40
-0.02461195  0.10758431  0.05624488  0.13749335

>

> ## Code for lasso with linear ##
> library(glmnet)
> lasso.lin.cv.errors = rep(NA,k)
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   x=model.matrix(Class~.,train)[,-1] #remove the int
ercept (training)
+   y=train[,dim(divorce)[2]] #response Class (training
)
+   cv.out=cv.glmnet(x,y,alpha=1)
+   bestlam=cv.out$lambda.min #get optimal tuning param
eter (lambda)
+
+   lasso.mod=glmnet(x,y,alpha=1,lambda=bestlam)
+   newx = model.matrix(Class~.,test )[-1] #remove the
intercept (test)
+   newy = test[,dim(divorce)[2]] #response Class (test
)
+   pred.lasso = predict(lasso.mod, s = bestlam, newx =
newx) #predict
+
+   for(m in 1:length(pred.lasso)){
+     if(pred.lasso[m]>=0.5){
+       pred.lasso[m]=1
+     } else{
+       pred.lasso[m]=0

```



```

+     }
+ }
+
+ #find MSE of lasso on this kth fold
+ error = mean((newy - pred.lasso)^2)
+ error
+
+ #append
+ lasso.lin.cv.errors[j] = error
+
+ #to see the coefs of the kth-fold's lowest MSE
+ lasso.coef=coef(lasso.mod)[,1]
+ print(lasso.coef[lasso.coef!=0])
+
+ }

```

(Intercept)	Atr2	Atr4	
-6.398876e-02	2.944569e-05	3.606098e-03	
Atr6	Atr7	Atr12	
5.682331e-02	2.682467e-02	2.666280e-03	
Atr15	Atr17	Atr18	
4.045280e-02	3.960235e-02	4.058390e-02	
Atr19	Atr26	Atr28	
2.600655e-02	3.248353e-02	1.552487e-02	
Atr31	Atr33	Atr36	
7.258096e-03	5.864019e-03	7.877197e-03	
Atr38	Atr40	Atr46	
1.350457e-02	5.623974e-02	-3.187309e-03	
Atr49	Atr52		
9.024248e-03	1.183193e-03		
(Intercept)	Atr3	Atr6	Atr11
-0.053560046	0.010992800	0.046967209	0.014623062
Atr17	Atr18	Atr26	Atr28
0.023043379	0.059092640	0.027553308	0.024959684
Atr31	Atr40	Atr49	Atr52
0.002165845	0.120198830	0.011739784	0.003335036
(Intercept)	Atr1	Atr2	Atr3
-0.090281679	0.010927643	0.009827424	0.016670085
Atr6	Atr7	Atr9	Atr11
0.048197398	0.029782146	0.014645164	0.007879829
Atr18	Atr19	Atr28	Atr29
0.057756373	0.022514579	0.018123875	0.014030075
Atr32	Atr39	Atr40	Atr49
0.002612150	0.022668912	0.077307020	0.009077201
Atr52			
0.016467939			
(Intercept)	Atr2	Atr3	
-0.0700811666	0.0054913956	0.0103762518	
Atr6	Atr7	Atr11	
0.0605669369	0.0099366964	0.0093607290	
Atr15	Atr17	Atr18	
0.0263838139	0.0316957435	0.0547311309	
Atr26	Atr28	Atr40	
0.0236364301	0.0076000893	0.1015792432	
Atr46	Atr48	Atr49	
-0.0079609062	-0.0006996802	0.0189200665	

```

      Atr52
0.0164566133
(Intercept)      Atr3      Atr6      Atr7
-0.069396665  0.010130010  0.047650860  0.015954611
      Atr9      Atr11      Atr17      Atr18
0.017736409  0.032499411  0.020232770  0.048387391
      Atr26      Atr30      Atr40      Atr41
0.017096396  0.004577569  0.087597236  0.022358047
      Atr44      Atr49      Atr52
0.019212873  0.005078399  0.009978129
> lasso.lin.cv.errors
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> ## Code for lasso with logistic##
> lasso.log.cv.errors = rep(NA,k)
> for(j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce[folds ==j,]
+
+   x=model.matrix(Class~.,train)[,-1] #remove the intercept (training)
+   y=train[,dim(divorce)[2]] #response Class (training)
+
+   cv.out=cv.glmnet(x,y,alpha=1)
+   bestlam=cv.out$lambda.min #get optimal tuning parameter (lambda)
+
+   lasso.mod=glmnet(x,y,alpha=1,lambda=bestlam, family="binomial")
+   newx = model.matrix(Class~.,test)[,-1] #remove the intercept (test)
+   newy = test[,dim(divorce)[2]] #response Class (test)
+
+   pred.lasso = predict(lasso.mod, s = bestlam, newx = newx, type = "response") #predict
+   pred.lasso[pred.lasso>=.5]=1
+   pred.lasso[pred.lasso<.5]=0
+   print(pred.lasso)
+   print(newy)
+
+   #find MSE of lasso on this kth fold
+   error = mean((newy != pred.lasso)^2)
+   error
+
+   #append
+   lasso.log.cv.errors[j] = error
+
+   #to see the coefs of the kth-fold's lowest MSE
+   lasso.coef=coef(lasso.mod)[,1]
+   print(lasso.coef[lasso.coef!=0])
+
+
+ }
1
1 0

```

3 1
 10 0
 15 1
 18 1
 19 1
 24 1
 26 1
 36 1
 37 1
 39 1
 46 1
 51 1
 54 1
 55 1
 63 1
 66 1
 67 1
 78 1
 89 0
 94 0
 98 0
 101 0
 107 0
 108 0
 111 0
 115 0
 121 0
 123 0
 131 0
 139 0
 146 0
 154 0
 156 0
 157 0
 160 0
 165 0
 166 0
 169 0

[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0

[25] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0

(Intercept)	Atr6	Atr15	Atr17
-5.8557004	0.5374038	0.6329337	0.2120006
Atr19	Atr26	Atr31	Atr40
0.7012349	1.0152343	0.1373878	0.7805287
Atr49			
0.1622556			

1
 4 1
 7 1
 13 1
 14 1
 22 1
 23 1
 29 1
 30 1
 34 1

```

40 1
42 1
43 1
45 1
69 1
70 1
72 1
76 1
77 1
81 1
82 1
84 1
91 0
110 0
112 0
120 0
125 0
127 0
128 0
129 0
132 0
140 0
142 0
143 0
153 0
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0
[25] 0 0 0 0 0 0 0 0 0 0 0 0
(Intercept)      Atr3      Atr6      Atr18
-5.1470283    0.1582206    0.4146123    0.3008922
      Atr26      Atr28      Atr40      Atr49
0.6989765    0.1289404    1.4401921    0.2986299
1
6 0
8 1
9 1
28 1
41 1
47 1
48 1
50 1
68 1
71 1
74 1
79 1
80 1
90 0
92 0
93 0
102 0
103 0
104 0
105 0
113 0
116 0
118 0
130 0

```

```

133 0
134 0
135 0
136 0
138 0
155 0
161 0
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0
[25] 0 0 0 0 0 0 0 0
(Intercept)      Atr1      Atr2      Atr3
-5.7155865      0.1406552    0.1078151    0.2189390
      Atr6      Atr18      Atr19      Atr28
  0.4251436      0.3822256    0.4000094    0.2659912
      Atr39      Atr40      Atr49      Atr52
  0.3082291      0.9439608    0.2677671    0.1953335
1
2  1
25 1
27 1
31 1
32 1
33 1
35 1
38 1
49 1
52 1
56 1
59 1
61 1
62 1
73 1
86 0
88 0
96 0
97 0
106 0
109 0
114 0
119 0
122 0
124 0
141 0
144 0
150 0
151 0
152 0
158 0
162 0
164 0
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0
[25] 0 0 0 0 0 0 0 0 0 0 0
(Intercept)      Atr3      Atr6      Atr15
-6.18616911      0.27657522    0.64870110    0.05288914
      Atr18      Atr19      Atr26      Atr40
  0.28630154      0.10886982    0.85841547    1.46415897
      Atr49      Atr52

```

```

0.46538645 0.13650625
1
5 0
11 1
12 1
16 1
17 1
20 1
21 1
44 1
53 1
57 1
58 1
60 1
64 1
65 1
75 1
83 1
85 0
87 0
95 0
99 0
100 0
117 0
126 0
137 0
145 0
147 0
148 0
149 0
159 0
163 0
167 0
168 0
170 0
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0
[25] 0 0 0 0 0 0 0 0 0 0
(Intercept)      Atr3      Atr6      Atr11
-5.36235545  0.30329654  0.37871943  0.29916399
      Atr17      Atr18      Atr26      Atr40
 0.24353393  0.24338547  0.34377845  1.20978857
      Atr44      Atr52
 0.42292508  0.08235102
> lasso.log.cv.errors
[1] 0.05128205 0.00000000 0.03225806 0.00000000
[5] 0.03030303
> # Linear regression with all 54 parameters - RSS
> mean(lin.mod.mse_vector)
[1] 0.03678986
> # Linear regression with all 54 parameters - MSE / Class Error
> mean(lin.mod.class.mse_vector)
[1] 0.02276863
> # Polynomial regression (exponent of 2) with all 54 parameters - RSS
> mean(log.mod.mse_vector)

```

```

[1] 0.01566936
> # Polynomial regression (exponent of 2) with all 54 p
arameters - MSE / Class Error
> mean(log.mod.class.mse_vector)
[1] 0.01764042
> # Polynomial logistic regression (exponent of 2) with
all 54 parameters
> mean(log.mod.mse_vector1)
[1] 0.01566936
> # LDA with all 54 parameters
> mean(lda.error)
[1] 0.02276863
> # KNN with all 54 parameters (CV for value of K in 5-
fold)
> mean(knn.error)
[1] 0.02276863
> # Decision tree with all 54 parameters
> mean(tree.error)
[1] 0.02882924
> # LDA with only top 5 correlated parameters
> mean(lda5.error)
[1] 0.02276863
> # QDA with only top 5 correlated parameters
> mean(qda5.error)
[1] 0.01764042
> # KNN with only top 5 correlated parameters (CV for v
alue of k in 5-fold)
> mean(knn5.error)
[1] 0.3243866
> # Decision tree with only top 5 correlated parameters
> mean(tree5.error)
[1] 0.03030303
> # Logistic regression with all 54 parameters
> mean(cv.logi_error1.5)
[1] 0.0387848
> # Logistic regression (with control) with all 54 para
meters
> mean(cv.logiWC_error1.5)
[1] 0.03148009
> # GAM w/ Smoothing Splines and 5 df
> mean(gam.error)
[1] 0.02865098
> # Linear regression with best subset selection
> mean(mean.best.lin.cv.errors)
[1] 0.02330765
> # Linear regression with forward selection
> mean(mean.fwd.lin.cv.errors)
[1] 0.02073065
> # Linear regression with backward selection
> mean(mean.bwd.lin.cv.errors)
[1] 0.02320042
> # Linear regression with lasso
> mean(lasso.lin.cv.errors)
[1] 0.02276863
> # Logistic regression with lasso
> mean(lasso.log.cv.errors)

```

```
[1] 0.02276863
```

```
> divorce_test <- read.csv('C:/Users/14014/Documents/Cornell_Fall_2019/STSCI_4740/STSCI_4740_FinalProject/divorce_fun_test.csv', header=TRUE)
> attach(divorce_test)
```

```
The following objects are masked from divorce (pos = 4)
:
```

```
Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9
```

```
The following objects are masked from divorce (pos = 5)
:
```

```
Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9
```

```
The following objects are masked from divorce (pos = 6)
:
```

```
Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9
```

```
The following objects are masked from divorce (pos = 16)
):
```

```
Atr1, Atr10, Atr11, Atr12, Atr13, Atr14,
Atr15, Atr16, Atr17, Atr18, Atr19, Atr2,
Atr20, Atr21, Atr22, Atr23, Atr24, Atr25,
Atr26, Atr27, Atr28, Atr29, Atr3, Atr30,
Atr31, Atr32, Atr33, Atr34, Atr35, Atr36,
Atr37, Atr38, Atr39, Atr4, Atr40, Atr41,
Atr42, Atr43, Atr44, Atr45, Atr46, Atr47,
Atr48, Atr49, Atr5, Atr50, Atr51, Atr52,
Atr53, Atr54, Atr6, Atr7, Atr8, Atr9
```



```

> # With QDA and 5 parameters only
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce_test
+
+   qda.fit=qda(as.factor(Class)~Atr40+Atr17+Atr19+Atr1
8+Atr11, data=train)
+   qda.class=predict(qda.fit,test)$class
+   print(qda.class)
+ }
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
> # With Decision Tree with only 5 parameters
> for (j in 1:k){
+   set.seed(1)
+   train = divorce[folds!=j,]
+   test = divorce_test
+
+   tree5.divorce=tree(as.factor(Class)~.,train)
+   tree5.pred=predict(tree5.divorce,test,type="class")
+   print(tree5.pred)
+ }
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
[1] 0 0
Levels: 0 1
>
>

```